

AI for Human Resources

Hire fairer · understand your data · scale support safely

Where AI fits in HR

Recruit

Draft posts, rubrics,
fair questions

Understand

Read files, forecast,
analyse surveys

Scale

Agents for
repeat questions

AI takes the routine work so you focus on people. Adoption in HR jumped from ~26% to ~43% in a year.

Faster and fairer hiring

Hiring: Before AI vs With AI

BEFORE AI — manual & inconsistent

Job posts written from scratch

Résumés piled up, skimmed by hand

Interview quality depends on who asks

Bias easy to miss, hard to track

WITH AI — faster & fairer

Drafts a job post in minutes

Scans & groups CVs by skills in seconds


Shared rubric, consistent questions

Flags biased wording for review

AI clears routine work so HR can focus on the people behind the documents.

AI clears routine work; your judgment keeps it fair.

From job post to fair interview



Grounding: answer from YOUR documents

Give AI approved files (handbook, policies, surveys) so answers come from your data — not generic guesses.

Input	Approved docs, spreadsheets, contracts
AI role	Summarise, compare, flag risk
Output	Grounded, checkable answers

The guardrails that matter

- AI learns from the past — bias can carry forward
- Never upload PII to unapproved tools
- GDPR / CCPA still apply
- Anonymise while exploring

Scrapped

Amazon's biased résumé tool, after it penalised
“women's”.

Prompt vs Agent

Prompt vs Agent

PROMPT — a conversation	AGENT — a service desk
Starts fresh every time	Rules set once, reused always
You give context each time	Grounded in approved sources
Great for one-off, changing tasks	Consistent answers at scale
You review every output	Knows when to hand off to a human

Agents aren't more powerful than prompts — just more consistent and scoped.

Prompt = a conversation. Agent = a service desk. Use an agent when questions repeat and consistency matters.

Designing a safe HR agent

The 4 Building Blocks of a Safe HR Agent



In AI + HR, narrower is usually safer.

In AI + HR, narrower is usually safer.

Key Takeaways

1

AI drafts and surfaces — humans check for bias and decide.

2

Ground AI in approved documents; protect PII under GDPR/CCPA.

3

Prompt for one-off work; agent for repeat, consistency-critical work.

4

Agents need maintenance — a monthly review catches drift.

Questions?